

Predicting Bitcoin Price Movement Utilizing Individual Stock Prices

Paola Rodriguez, Tommy Barron

University of San Diego

ADS-502-01: Applied Data Mining

An Tran

December 9, 2024

Abstract

The increasing volatility of cryptocurrency markets, particularly Bitcoin, has created a demand for predictive models that can provide insights into price fluctuations. This study investigates the potential of using individual stock prices as predictors for Bitcoin price movements, leveraging machine learning models to analyze patterns and relationships. The dataset includes daily stock prices from major companies such as Apple, Tesla, Microsoft, Google, Nvidia, Berkshire Hathaway, Netflix, Amazon, and Meta, alongside Bitcoin's historical price data. We conducted exploratory data analysis (EDA) to identify trends, correlations, and distribution characteristics within the data. Subsequently, we implemented and evaluated multiple predictive models, including Naive Bayes, random forests, and neural networks, to determine their effectiveness in forecasting Bitcoin price changes. Additionally, lagged variables were incorporated to capture temporal dependencies in Bitcoin's price movement. The results indicate varying levels of predictive accuracy across the models, highlighting the challenges and opportunities in predicting Bitcoin price movements using stock market data. This paper contributes to the understanding of cross-market influences and proposes a framework for leveraging machine learning in financial prediction. The findings have implications for traders, investors, and researchers aiming to enhance their strategies in cryptocurrency markets.

Keywords: cryptocurrency markets, individual stock prices, predictive models, lagged variables

Table of Contents

Abstract	2
Table of Contents	3
Methodology	4
C5.0	5
CART	6
Logistic Regression	8
Random Forests	10
Naïve Bayes	10
Neural Networks	11
Results	12
Conclusion	14
References	15
Appendix	16

Predicting Bitcoin Price Movement Utilizing Individual Stock Prices

The volatility of Bitcoin and its rapid growth as a digital asset have made it a focal point for predictive modeling. Understanding and anticipating its price movements is of significant interest to traders, analysts, and researchers. While various economic indicators and cryptocurrency-specific metrics already exist and have been used to forecast Bitcoin prices, this study explores the predictive potential of stock market data.

This paper examines whether individual stock prices, including those of major companies such as Apple, Tesla, and Microsoft, can be leveraged to forecast Bitcoin's daily price movements. By applying machine learning techniques, including neural networks, linear regression, and Naive Bayes, this analysis evaluates the relationship between stock market trends and Bitcoin price changes. The findings aim to contribute to financial modeling strategies and expand the understanding of cross-market correlations.

Methodology

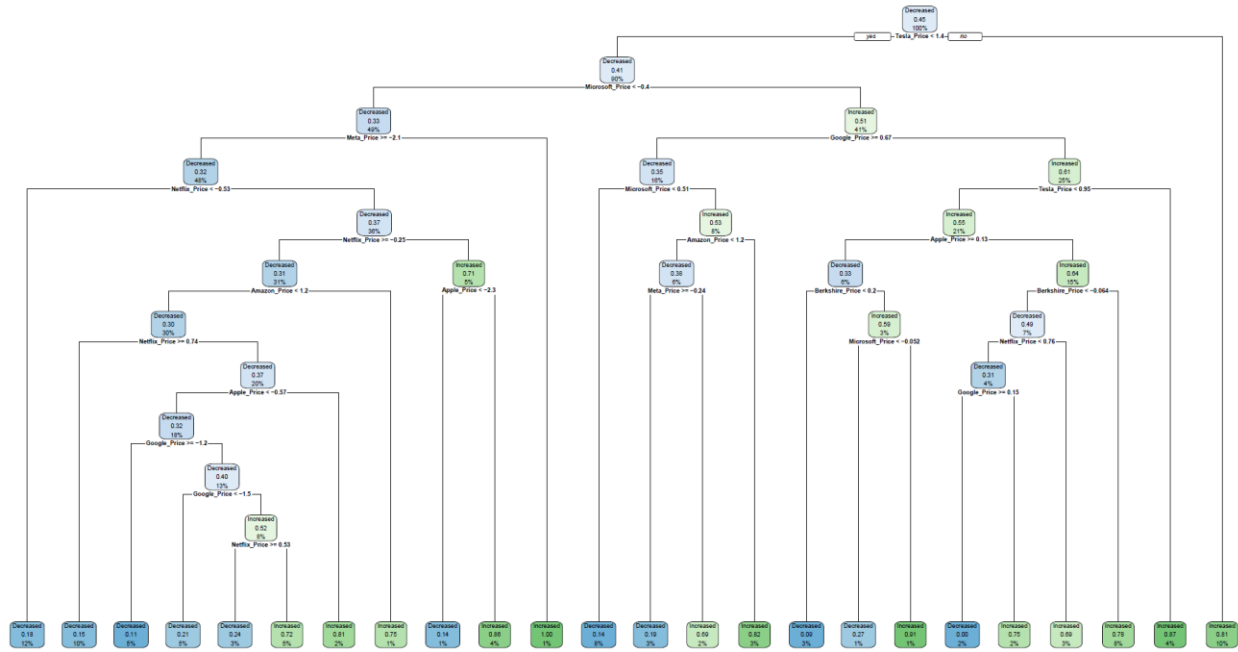
The data that was chosen for this project is a collection of US stock market and commodities data from 2020-2024. The data contains 37 columns of numerical data, a date column, and an index column. The dates were not formatted correctly and were standardized in the YYYY-MM-DD format. The data was then split into a training and test data set where 75% of the data is in the training set and the rest is in the test set. The data was split exactly at 2023-01-31. It was determined that a date split is preferable instead of random sampling because in the domain of the stock market, a data scientist cannot use future data to predict future prices. Since there are no missing variables, the test and train numerical data were then standardized to assist with the models. For this project, the models will be focusing on using individual stock prices to predict Bitcoin price movement. To achieve this goal, Bitcoin prices were lagged by a day so the

current stock prices can predict the predictor variable in the models. A rolling seven-day average for Bitcoin prices was calculated for each instance and each instance was compared to the last day; therefore, if the price increases, the movement column will indicate an increase and vice versa for decreasing prices. This creates the binary predictor variable of Bitcoin price movement where increases are 1 and decreases are 0. Finally, the train and test data were inspected for balancing issues and was confirmed to be balanced at 412 decreases and 341 increases for the train set and 133 decreases and 113 increases for the test set.

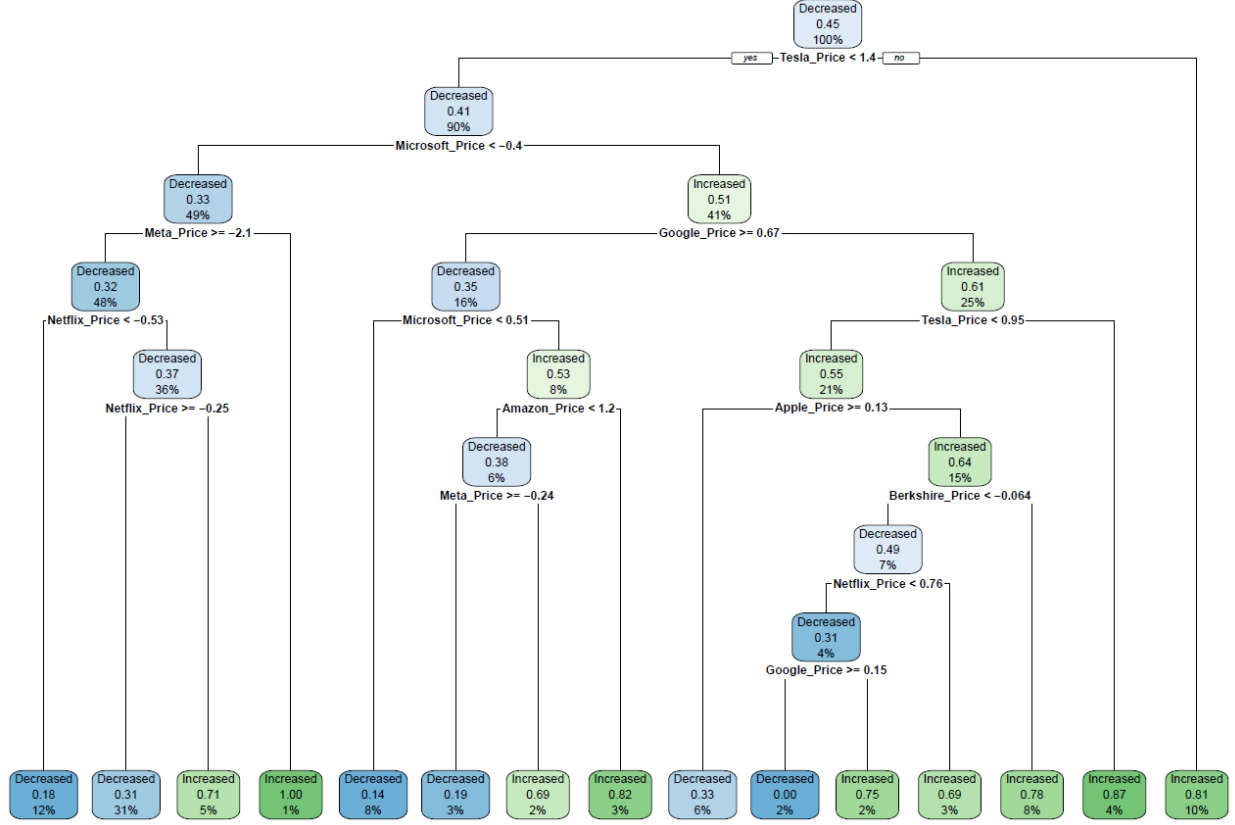
C5.0

The C5.0 decision tree highlights the relationship between individual stock prices and Bitcoin's price movement, identifying variables like Nvidia_Price, Amazon_Price, and Apple_Price as significant predictors. The tree's structure prioritizes Nvidia_Price as the initial decision node, emphasizing its importance in determining Bitcoin's direction. However, the model exhibits a bias toward predicting "Increase," as shown in the confusion matrix, where most "Decrease" instances are misclassified as "Increase." While the tree provides detailed decision boundaries based on stock price thresholds, its depth and complexity may lead to overfitting, reducing generalizability.

To improve the model, steps such as pruning the tree, balancing the training data, and performing cross-validation can address bias and enhance interpretability. Additionally, simplifying the tree and analyzing feature importance could help refine the predictors further. Overall, the model demonstrates potential in using individual stock prices to predict Bitcoin movements but requires optimization for better performance and reliability.

Figure 2*CART Decision Tree for Predicting Bitcoin Price Movement*

The complex nature of the decision tree created indicates that there are some issues with overfitting due to high variance (Geeks for Geeks, 2024). Additional pruning and regularization of parameters will need to occur to limit these issues. Figure 3 is a second CART model that was developed using pruning tactics to understand the impact of certain variables. The model was pruned with a complexity parameter of 0.015 and has a classification threshold set at 0.3. The resulting decision tree has 15 leaves and 14 decision nodes with overall increased model evaluation metrics, most notably the F-scores.

Figure 3*Pruned CART Decision Tree for Predicting Bitcoin Price Movement***Logistic Regression**

Between linear and logistic regression, logistic regression was chosen to best represent the data because of the binary target variable. The generalized logistic regression equation takes the form:

$$p(y) = \frac{\exp(\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p)}{1 + \exp(\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p)} + \varepsilon \quad (1)$$

Where β_p represents predictor variable coefficients and x_p represents the predictor variable value. Therefore, plugging in the values calculated from the logistic regression function yields:

$$\hat{p}(\text{Movement}) = \frac{\exp(0.3216 - 1.8567(\text{Apple}) + 0.8394(\text{Tesla}) + 1.6823(\text{Microsoft}) - 1.4646(\text{Google}) + 0.8578(\text{Nvidia}) + 0.8860(\text{Berkshire}) - 0.6075(\text{Netflix}) + 0.2561(\text{Amazon}) + 0.4536(\text{Meta}))}{1 + \exp(0.3216 - 1.8567(\text{Apple}) + 0.8394(\text{Tesla}) + 1.6823(\text{Microsoft}) - 1.4646(\text{Google}) + 0.8578(\text{Nvidia}) + 0.8860(\text{Berkshire}) - 0.6075(\text{Netflix}) + 0.2561(\text{Amazon}) + 0.4536(\text{Meta}))} \quad (2)$$

It is important to note that the logistic regression has calculated varying p-values for the predictor variables. The variables that showed the least significance at $\alpha > 0.05$ are Nvidia, Amazon, and Meta. Therefore, the variables that showed the greatest significance are Apple, Tesla, Microsoft, Google, Berkshire, and Netflix. When evaluating the regression model, the area under the receiver operating characteristic curve was calculated to measure the model's discrimination rate. The area under the curve came out to be 0.505 which indicates that this model is no better than random guessing.

A second logistic regression model was created after removing the statistically insignificant predictor variables (p-value > 0.05). While creating the model, the variables that were deemed unimpactful were Netflix, Amazon, Nvidia, and Meta. After these variables were removed the resulting area under the curve to be 0.5289. This indicates a “stronger” model but is still quite low and may need further development. The formula for the second model is then made to be:

$$\hat{p}(\text{Movement}) = \frac{\exp(0.3216 - 1.8567(\text{Apple}) + 0.8394(\text{Tesla}) + 1.6823(\text{Microsoft}) - 1.4646(\text{Google}) + 0.8860(\text{Berkshire}))}{1 + \exp(0.3216 - 1.8567(\text{Apple}) + 0.8394(\text{Tesla}) + 1.6823(\text{Microsoft}) - 1.4646(\text{Google}) + 0.8860(\text{Berkshire}))} \quad (3)$$

Random Forests

The goal of the random forest model is to combine multiple decision trees together into a single output (IBM, n.d.). With a relatively larger amount of predictor variables, the random forests model generated 500 trees. The model produced an accuracy of 54.07%, a sensitivity of 81.42%, a specificity of 30.83, a precision of 44.57%, and an AUC-ROC of 0.5612. During the development, it was speculated that adjusting the number of trees may improve the model. Two additional models were developed to test against the baseline. The first model had an “ntree” count of 1000 trees. When evaluating the model, only specificity went up from 0.3083 to 0.4812, but this comes at the cost of decreasing sensitivity from 0.8142 to 0.6460. The slight increase in specificity does not justify the drastic decrease in the overall metrics. The second model was developed using 100 trees. This model’s performance decreased with the three-evaluation metrics, sensitivity, specificity, and recall. Therefore, the baseline, 500 trees, is the preferred random forest model.

Naïve Bayes

The Naive Bayes model for predicting Bitcoin price movement based on individual stock prices demonstrates moderate accuracy at 50.42%. From the confusion matrix, it is evident that the model performs better at predicting "Increase" (85 correct predictions) compared to "Decrease" (35 correct predictions). However, the model struggles with misclassification, as 93 "Actual: Decrease" instances are predicted as "Increase," and 25 "Actual: Increase" instances are incorrectly classified as "Decrease."

Our basis for this approach hinges on Bayes Theorem:

$$p(X *) = \frac{p(Y = y *)p(Y = y *)}{p(X *)} \quad (4)$$

where x and y are the posterior probabilities (Larose, 2019, pp.113-114). This performance indicates that the Naive Bayes model has limited predictive power in this context, likely due to its assumption of feature independence, which may not hold true for stock prices that are often correlated. To improve accuracy, feature engineering, reducing noise in the data, or incorporating a different algorithm better suited for capturing complex relationships, such as Random Forest or Gradient Boosting, could enhance the predictive performance. Nonetheless, the model provides a baseline for understanding the relationship between stock prices and Bitcoin price movement.

Neural Networks


The neural network model developed for this research analyzes the relationship between individual stock prices and Bitcoin price movement. The model consists of an input layer representing nine major stock prices (Apple, Tesla, Microsoft, Google, Nvidia, Berkshire Hathaway, Netflix, Amazon, and Meta), a hidden layer for processing these inputs, and an output layer predicting Bitcoin price movement. The connections between the input and hidden layers are weighted, reflecting the relative importance of each stock in influencing the model's predictions.

The results indicate that Nvidia and Berkshire Hathaway stock prices have the highest weights, suggesting they are the most influential in predicting Bitcoin price changes. This finding highlights potential stronger correlations between these specific stocks and Bitcoin's movement compared to others like Google or Meta, which have comparatively lower weights. By processing the non-linear relationships between these inputs, the model effectively captures complex dependencies, providing a nuanced understanding of how traditional equity markets

interact with cryptocurrency prices. These insights could prove valuable for developing more robust predictive models in financial markets.

Figure 2

Neural Network Predicting Bitcoin Price



Results

The results shown in Table 1 display the model evaluation metrics shown for the models described in the methodology. To choose the best model, it is necessary for the model to perform well overall because predicting increases in Bitcoin movement is just as important as predicting decreases. The logistic regression model has a few well performing metrics of the F_1 -score and the F_2 -score at 61.92% and 75.53%, respectively. The relatively high F_1 -score indicates the model has a good balance between precision and recall while the high F_2 -score shows a particular strength with recall (Wood, n.d.). The C5.0 model has three strong metrics with the greatest sensitivity at 90.91%, specificity at 62.50%, and F_2 -score at 75.76%. This implies that the C5.0 model performs well at classifying positive and negative records with a particular emphasis on recall rate. Although these metrics are relatively strong, the C5.0 performs the

weakest with the highest error rate of 54.62% and one of the lowest precision scores of 45.45%.

To choose a well-rounded model, it is important to have strong metrics, but also to limit the weakest metrics. Therefore, the best model to consider for this project is the CART decision tree. This model has the highest accuracy at 56.10%, precision at 52.03%, $F_{0.5}$ -score at 52.89, and the greatest ROC-AUC at 56.14%. The high accuracy indicates that it has the highest proportion of correct to incorrect classifications. The $F_{0.5}$ -score and, subsequently, precision demonstrates the model's ability to properly classify positives as true positives rather than false positives.

Additionally, the greatest ROC-AUC indicates the highest proportion of actual positives and negatives identified from the data. Although the CART model has the lowest recall rate, it is important to note that the other models are not well-rounded and lack significantly in specificity and/or precision. Since both identifying increases and decreases in Bitcoin price movement is essential, the preferred method is the CART decision tree.

Table 1

Model Evaluation Table for Bitcoin Price Movement

Evaluation Measure	C5.0	CART	Logistic Regression	Random Forest	Naïve Bayes	Neural Network
Accuracy	0.4538	0.5610	0.5000	0.5407	0.5042	0.4664
Error Rate	0.5462	0.4390	0.5000	0.4593	0.5000	0.5340
Sensitivity	0.9091	0.5664	0.8850	0.8142	0.7727	0.7181
Specificity	0.6250	0.5564	0.1729	0.3083	0.2734	0.2500
Precision	0.4545	0.5203	0.4762	0.4457	0.4780	0.4514
F_1	0.5884	0.5424	0.6192	0.5761	0.5906	0.5420
F_2	0.7576	0.5565	0.7553	0.6987	0.6878	0.6422
$F_{0.5}$	0.5049	0.5289	0.5247	0.4901	0.5174	0.2194
AUC	0.5137	0.5614	0.5289	0.5612	0.4912	0.4513

Conclusion

From the models developed in this project, the recommended model for predicting Bitcoin price movement is the CART decision tree due to its best overall positive and negative identification rates. Although this was the best performing model, it is important to consider the limitations. Most notably, within the scope of this project predicting the stock market and cryptocurrency is almost an impossible task due to the complexity and deeper influences of product development and the economy at large. Politics, public relations, and other non-numeric factors move prices in ways that are more complicated than simply using individual stocks as predictor variables. Thus, more research must be done to find any useful implications to real-world applications. Future developments using current world affairs, politics, and other categorical variables may prove useful for domain expansion across the stock market.

References

- Geeks for Geeks. (2024, Sep 19). Geeks for Geeks. CART (Classification And Regression Tree) in Machine Learning. Geeks for Geeks. <https://www.geeksforgeeks.org/cart-classification-and-regression-tree-in-machine-learning/>
- IBM. (n.d.). What is Random Forest? IBM. <https://www.ibm.com/topics/random-forest>
- Larose, C., & Larose, D. (2019). Data science using Python and R. John Wiley & Sons, Inc.
- Tan, P. N., Steinbach, M., Karpatne, A., & Kumar, V. (2019). Introduction to data mining (2nd ed.). Pearson.

Appendix

Upload data and check for N/A values

```
df <- read.csv("us_stock_data.csv")
head(df)
```

```
##      X      Date Natural_Gas_Price Natural_Gas_Vol. Crude_oil_Price
## 1 0    2/2/2024          2.079              NA          72.28
## 2 1    1/2/2024          2.050             161340          73.82
## 3 2 31-01-2024          2.100             142860          75.85
## 4 3 30-01-2024          2.077             139750          77.82
## 5 4 29-01-2024          2.490              3590          76.78
## 6 5 26-01-2024          2.712             73020          78.01
##      Crude_oil_Vol. Copper_Price Copper_Vol. Bitcoin_Price Bitcoin_Vol.
## 1              NA       3.8215              NA    43,194.70    42650
## 2          577940       3.8535              NA    43,081.40    47690
## 3          344490       3.9060              NA    42,580.50    56480
## 4          347240       3.9110              NA    42,946.20    55130
## 5          331930       3.8790              NA    43,299.80    45230
## 6          365460       3.8520              NA    41,811.30    69470
##      Platinum_Price Platinum_Vol. Ethereum_Price Ethereum_Vol. S.P_500_Price
## 1           901.6          NA       2,309.28       246890    4,958.61
## 2           922.3          NA       2,304.28       323610    4,906.19
## 3           932.6          NA       2,283.14       408790    4,848.87
## 4           931.7          NA       2,343.11       387120    4,924.97
## 5           938.3          NA       2,317.79       318840    4,927.93
## 6           921.3          NA       2,267.55       377790    4,890.97
##      Nasdaq_100_Price Nasdaq_100_Vol. Apple_Price Apple_Vol. Tesla_Price
## 1          17,642.73      315620000       185.85  102550000    187.91
## 2          17,344.71      240640000       186.86   53490000    188.86
## 3          17,137.24      366450000       184.40   54830000    187.29
## 4          17,476.71      236210000       188.04   55270000    191.59
## 5          17,596.27      238750000       191.73   46890000    190.93
## 6          17,421.01      252940000       192.42   44590000    183.25
##      Tesla_Vol. Microsoft_Price Microsoft_Vol. Silver_Price Silver_Vol.
## 1  110610000       411.22      28260000       22.796          NA
## 2   90680000       403.78      29230000       23.236       85160
## 3  102270000       397.58      46780000       23.169       66910
## 4  105540000       408.59      29340000       23.225       53370
## 5  123600000       409.72      23290000       23.134         330
## 6  107340000       403.93      17800000       22.758         330
##      Google_Price Google_Vol. Nvidia_Price Nvidia_Vol. Berkshire_Price
## 1          142.38    62500000       661.60    47660000      5,89,498
## 2          141.16    37120000       630.27    36020000      5,81,600
## 3          140.10    71370000       615.27    45070000      5,78,020
## 4          151.46    33060000       627.74    39600000      5,84,680
## 5          153.51    27590000       624.65    33900000      5,78,800
```



```
## 6      152.18      26120000      610.31      39030000      5,82,300
##  Berkshire_Vol. Netflix_Price Netflix_Vol. Amazon_Price Amazon_Vol. Meta_
Price
## 1      10580      564.64      4030000      171.81      117220000      4
74.99
## 2      9780      567.51      3150000      159.28      66360000      3
94.78
## 3      9720      564.11      4830000      155.20      49690000      3
90.14
## 4      9750      562.85      6120000      159.00      42290000      4
00.06
## 5      13850      575.79      6880000      161.26      42840000      4
01.02
## 6      10040      570.42      12770000      159.12      51050000      3
94.14
##  Meta_Vol. Gold_Price Gold_Vol.
## 1 84710000 2,053.70      NA
## 2 25140000 2,071.10      260920
## 3 20010000 2,067.40      238370
## 4 18610000 2,050.90      214590
## 5 17790000 2,034.90      1780
## 6 13160000 2,026.60      410
```

```
colSums(is.na(df))
```

```
##           X           Date Natural_Gas_Price Natural_Gas_Vol.
##           0           0           0           4
##  Crude_oil_Price Crude_oil_Vol.   Copper_Price   Copper_Vol.
##           0           23           0           37
##   Bitcoin_Price   Bitcoin_Vol.   Platinum_Price   Platinum_Vol.
##           0           0           0           377
##   Ethereum_Price   Ethereum_Vol.   S.P_500_Price   Nasdaq_100_Price
##           0           0           0           0
##  Nasdaq_100_Vol.   Apple_Price   Apple_Vol.   Tesla_Price
##           1           0           0           0
##   Tesla_Vol.   Microsoft_Price   Microsoft_Vol.   Silver_Price
##           0           0           0           0
##   Silver_Vol.   Google_Price   Google_Vol.   Nvidia_Price
##           46           0           0           0
##   Nvidia_Vol.   Berkshire_Price   Berkshire_Vol.   Netflix_Price
##           0           0           0           0
##   Netflix_Vol.   Amazon_Price   Amazon_Vol.   Meta_Price
##           0           0           0           0
##   Meta_Vol.   Gold_Price   Gold_Vol.
##           0           0           2
```

Data Pre-processing

Reformat date using lubridate package

```
df$standardized_dates <- parse_date_time(df$Date, orders = c("dmy"))
df$standardized_dates <- as.Date(df$standardized_dates)
head(df)
```

##	X	Date	Natural_Gas_Price	Natural_Gas_Vol.	Crude_oil_Price
## 1	0	2/2/2024	2.079	NA	72.28
## 2	1	1/2/2024	2.050	161340	73.82
## 3	2	31-01-2024	2.100	142860	75.85
## 4	3	30-01-2024	2.077	139750	77.82
## 5	4	29-01-2024	2.490	3590	76.78
## 6	5	26-01-2024	2.712	73020	78.01

##	Crude_oil_Vol.	Copper_Price	Copper_Vol.	Bitcoin_Price	Bitcoin_Vol.
## 1	NA	3.8215	NA	43,194.70	42650
## 2	577940	3.8535	NA	43,081.40	47690
## 3	344490	3.9060	NA	42,580.50	56480
## 4	347240	3.9110	NA	42,946.20	55130
## 5	331930	3.8790	NA	43,299.80	45230
## 6	365460	3.8520	NA	41,811.30	69470

##	Platinum_Price	Platinum_Vol.	Ethereum_Price	Ethereum_Vol.	S.P_500_Price
## 1	901.6	NA	2,309.28	246890	4,958.61
## 2	922.3	NA	2,304.28	323610	4,906.19
## 3	932.6	NA	2,283.14	408790	4,848.87
## 4	931.7	NA	2,343.11	387120	4,924.97
## 5	938.3	NA	2,317.79	318840	4,927.93
## 6	921.3	NA	2,267.55	377790	4,890.97

##	Nasdaq_100_Price	Nasdaq_100_Vol.	Apple_Price	Apple_Vol.	Tesla_Price
## 1	17,642.73	315620000	185.85	102550000	187.91
## 2	17,344.71	240640000	186.86	53490000	188.86
## 3	17,137.24	366450000	184.40	54830000	187.29
## 4	17,476.71	236210000	188.04	55270000	191.59
## 5	17,596.27	238750000	191.73	46890000	190.93
## 6	17,421.01	252940000	192.42	44590000	183.25

##	Tesla_Vol.	Microsoft_Price	Microsoft_Vol.	Silver_Price	Silver_Vol.
## 1	110610000	411.22	28260000	22.796	NA
## 2	90680000	403.78	29230000	23.236	85160
## 3	102270000	397.58	46780000	23.169	66910
## 4	105540000	408.59	29340000	23.225	53370
## 5	123600000	409.72	23290000	23.134	330
## 6	107340000	403.93	17800000	22.758	330

##	Google_Price	Google_Vol.	Nvidia_Price	Nvidia_Vol.	Berkshire_Price
## 1	142.38	62500000	661.60	47660000	5,89,498
## 2	141.16	37120000	630.27	36020000	5,81,600
## 3	140.10	71370000	615.27	45070000	5,78,020
## 4	151.46	33060000	627.74	39600000	5,84,680
## 5	153.51	27590000	624.65	33900000	5,78,800
## 6	152.18	26120000	610.31	39030000	5,82,300

```
## Berkshire_Vol. Netflix_Price Netflix_Vol. Amazon_Price Amazon_Vol. Meta_
Price
## 1          10580          564.64          4030000          171.81          117220000          4
74.99
## 2           9780          567.51          3150000          159.28          66360000          3
94.78
## 3           9720          564.11          4830000          155.20          49690000          3
90.14
## 4           9750          562.85          6120000          159.00          42290000          4
00.06
## 5          13850          575.79          6880000          161.26          42840000          4
01.02
## 6          10040          570.42          12770000          159.12          51050000          3
94.14
## Meta_Vol. Gold_Price Gold_Vol. standardized_dates
## 1 84710000 2,053.70          NA          2024-02-02
## 2 25140000 2,071.10        260920          2024-02-01
## 3 20010000 2,067.40        238370          2024-01-31
## 4 18610000 2,050.90        214590          2024-01-30
## 5 17790000 2,034.90         1780          2024-01-29
## 6 13160000 2,026.60         410          2024-01-26
```

We will be focusing on predicting Bitcoin Price Movement using individual stocks

Individual stocks: Apple, Tesla, Microsoft, Google, Nvidia, Berkshire Hathaway, Netflix, Amazon, and Meta

Create filtered data frame with individual stocks and dates

```
filtered_df <- df[, c("standardized_dates", "Bitcoin_Price", "Apple_Price", "
Tesla_Price", "Microsoft_Price", "Google_Price", "Nvidia_Price", "Berkshire_P
rice", "Netflix_Price", "Amazon_Price", "Meta_Price")]
```

```
filtered_df$Bitcoin_Price <- gsub("[^0-9.-]", "", filtered_df$Bitcoin_Price)
filtered_df$Bitcoin_Price <- as.numeric(filtered_df$Bitcoin_Price)
filtered_df$Berkshire_Price <- gsub("[^0-9.-]", "", filtered_df$Berkshire_Pri
ce)
filtered_df$Berkshire_Price <- as.numeric(filtered_df$Berkshire_Price)
head(filtered_df)
```

```
## standardized_dates Bitcoin_Price Apple_Price Tesla_Price Microsoft_Price
## 1          2024-02-02          43194.7          185.85          187.91          411.22
## 2          2024-02-01          43081.4          186.86          188.86          403.78
## 3          2024-01-31          42580.5          184.40          187.29          397.58
## 4          2024-01-30          42946.2          188.04          191.59          408.59
## 5          2024-01-29          43299.8          191.73          190.93          409.72
## 6          2024-01-26          41811.3          192.42          183.25          403.93
## Google_Price Nvidia_Price Berkshire_Price Netflix_Price Amazon_Price
## 1          142.38          661.60          589498          564.64          171.81
## 2          141.16          630.27          581600          567.51          159.28
```

```
## 3      140.10      615.27      578020      564.11      155.20
## 4      151.46      627.74      584680      562.85      159.00
## 5      153.51      624.65      578800      575.79      161.26
## 6      152.18      610.31      582300      570.42      159.12
## Meta_Price
## 1      474.99
## 2      394.78
## 3      390.14
## 4      400.06
## 5      401.02
## 6      394.14

# Sort the filtered data frame by date

sorted_filtered_df <- filtered_df[order(filtered_df$standardized_dates, decreasing = TRUE),]
head(sorted_filtered_df)

## standardized_dates Bitcoin_Price Apple_Price Tesla_Price Microsoft_Price
## 1      2024-02-02      43194.7      185.85      187.91      411.22
## 2      2024-02-01      43081.4      186.86      188.86      403.78
## 3      2024-01-31      42580.5      184.40      187.29      397.58
## 4      2024-01-30      42946.2      188.04      191.59      408.59
## 5      2024-01-29      43299.8      191.73      190.93      409.72
## 6      2024-01-26      41811.3      192.42      183.25      403.93
## Google_Price Nvidia_Price Berkshire_Price Netflix_Price Amazon_Price
## 1      142.38      661.60      589498      564.64      171.81
## 2      141.16      630.27      581600      567.51      159.28
## 3      140.10      615.27      578020      564.11      155.20
## 4      151.46      627.74      584680      562.85      159.00
## 5      153.51      624.65      578800      575.79      161.26
## 6      152.18      610.31      582300      570.42      159.12
## Meta_Price
## 1      474.99
## 2      394.78
## 3      390.14
## 4      400.06
## 5      401.02
## 6      394.14
```

Partition the Data into Test (25%) and Train (75%)

```
split_point <- floor(0.25 * nrow(sorted_filtered_df))

test_data <- sorted_filtered_df[1:split_point,]
train_data <- sorted_filtered_df[(split_point + 1):nrow(sorted_filtered_df),]

n_test <- dim(test_data)[1]
test_data$Index <- c(1:n_test)
```

```

n_train <- dim(train_data)[1]
train_data$Index <- c(1:n_train)

# Uncomment to check the test and train data

## Train data information:
## 760 Data points from 2020-01-02 to 2023-01-31

## Test data information:
## 253 Data points from 2023-02-01 to 2024-02-02

head(train_data)

##      standardized_dates Bitcoin_Price Apple_Price Tesla_Price Microsoft_Pri
ce
## 254      2023-01-31      23125.1      144.29      173.22      247.
81
## 255      2023-01-30      22832.2      143.00      166.66      242.
71
## 256      2023-01-27      23074.6      145.93      177.90      248.
16
## 257      2023-01-26      23016.0      143.96      160.27      248.
00
## 258      2023-01-25      23055.1      141.86      144.43      240.
61
## 259      2023-01-24      22632.5      142.53      143.89      242.
04
##      Google_Price Nvidia_Price Berkshire_Price Netflix_Price Amazon_Price
## 254      98.84      195.37      473000      353.86      103.13
## 255      96.94      191.62      465040      353.11      100.55
## 256      99.37      203.65      470000      360.77      102.24
## 257      97.52      198.02      469960      364.87      99.22
## 258      95.22      193.23      471658      367.96      97.18
## 259      97.70      192.65      471000      363.83      96.32
##      Meta_Price Index
## 254      148.97      1
## 255      147.06      2
## 256      151.74      3
## 257      147.30      4
## 258      141.50      5
## 259      143.14      6

head(test_data)

##      standardized_dates Bitcoin_Price Apple_Price Tesla_Price Microsoft_Price
## 1      2024-02-02      43194.7      185.85      187.91      411.22
## 2      2024-02-01      43081.4      186.86      188.86      403.78
## 3      2024-01-31      42580.5      184.40      187.29      397.58
## 4      2024-01-30      42946.2      188.04      191.59      408.59
## 5      2024-01-29      43299.8      191.73      190.93      409.72

```

```
## 6      2024-01-26      41811.3      192.42      183.25      403.93
##   Google_Price Nvidia_Price Berkshire_Price Netflix_Price Amazon_Price
## 1      142.38      661.60      589498      564.64      171.81
## 2      141.16      630.27      581600      567.51      159.28
## 3      140.10      615.27      578020      564.11      155.20
## 4      151.46      627.74      584680      562.85      159.00
## 5      153.51      624.65      578800      575.79      161.26
## 6      152.18      610.31      582300      570.42      159.12
##   Meta_Price Index
## 1      474.99      1
## 2      394.78      2
## 3      390.14      3
## 4      400.06      4
## 5      401.02      5
## 6      394.14      6
```

Create Test & Train sets that have standardized variables

Standardize the variables into a new data frame

```
std_sorted_filtered_df <- sorted_filtered_df
std_sorted_filtered_df[, 2:11] <- scale(sorted_filtered_df[, 2:11])
```

Split the new standardized data frame into the train and test sets

```
std_split_point <- floor(0.25 * nrow(std_sorted_filtered_df))

std_test <- std_sorted_filtered_df[1:std_split_point,]
std_train <- std_sorted_filtered_df[(std_split_point + 1):nrow(std_sorted_filtered_df),]
```

Create Index for the train and test sets

```
n_test <- dim(std_test)[1]
std_test$Index <- c(1:n_test)
n_train <- dim(std_train)[1]
std_train$Index <- c(1:n_train)
```

For time series analysis lag the target variable by one day to use previous day predictor variables to predict the movement of next day target variable

```
std_train$Bitcoin_Lag <- lag(std_train$Bitcoin_Price, 1)
std_test$Bitcoin_Lag <- lag(std_test$Bitcoin_Price, 1)
```

Create a seven day rolling mean for the target variable

```
std_train <- std_train %>%
  mutate(var1_7day_avg = rollmean(Bitcoin_Price, k=7, fill = NA, align = "right"))
```

```

std_test <- std_test %>%
  mutate(var2_7day_avg = rollmean(Bitcoin_Price, k=7, fill = NA, align = "right"))

## Use the seven day rolling mean to see if the predictor variable increases or decreases

std_train <- std_train %>%
  mutate(Bitcoin_Price_Change = ifelse(var1_7day_avg > lag(var1_7day_avg), "Increased",
                                     ifelse(var1_7day_avg < lag(var1_7day_avg), "Decreased", "No Change")))
std_test <- std_test %>%
  mutate(Bitcoin_Price_Change = ifelse(var2_7day_avg > lag(var2_7day_avg), "Increased",
                                     ifelse(var2_7day_avg < lag(var2_7day_avg), "Decreased", "No Change")))

std_train$Bitcoin_Price_Change <- factor(std_train$Bitcoin_Price_Change)
std_test$Bitcoin_Price_Change <- factor(std_test$Bitcoin_Price_Change)

std_train <- na.omit(std_train)
std_test <- na.omit(std_test)

head(std_train)

```

	standardized_dates	Bitcoin_Price	Apple_Price	Tesla_Price	Microsoft_Price
## 261	2023-01-20	-0.4325585	-0.1212140	-0.8891951	-0.44559
## 262	2023-01-19	-0.5381446	-0.1981854	-0.9627959	-0.58978
## 263	2023-01-18	-0.5653086	-0.1999617	-0.9438364	-0.52229
## 264	2023-01-17	-0.5344465	-0.1783505	-0.9119230	-0.44333
## 265	2023-01-13	-0.6145027	-0.2132836	-1.0189681	-0.46281
## 266	2023-01-12	-0.6856674	-0.2532495	-1.0053078	-0.47533

	Google_Price	Nvidia_Price	Berkshire_Price	Netflix_Price	Amazon_Price
## 261	-0.4182623	-0.3213529	0.5416961	-0.6640914	-1.470477
## 262	-0.6136369	-0.4048231	0.4485979	-0.8871933	-1.601141
## 263	-0.6895068	-0.3572591	0.5007425	-0.7991048	-1.535992
## 264	-0.6828239	-0.3320004	0.6185823	-0.8000232	-1.514397
## 265	-0.6501960	-0.3944088	0.6793290	-0.7449157	-1.438634
## 266	-0.6891137	-0.4245637	0.7015782	-0.7673762	-1.542946

	Meta_Price	Index	Bitcoin_Lag	var1_7day_avg	Bitcoin_Price_Change
## 261	-1.546021	8	-0.4167934	-0.4187346	Decreased

```
## 262 -1.590161      9 -0.4325585 -0.4352833      Decreased
## 263 -1.633068     10 -0.5381446 -0.4580034      Decreased
## 264 -1.600990     11 -0.5653086 -0.4757607      Decreased
## 265 -1.578783     12 -0.5344465 -0.5053243      Decreased
## 266 -1.582484     13 -0.6145027 -0.5410602      Decreased
```

```
head(std_test)
```

```
##      standardized_dates Bitcoin_Price Apple_Price Tesla_Price Microsoft_Pri
ce
## 8      2024-01-24      0.7191500      1.555282 -0.012933283      2.3780
22
## 9      2024-01-23      0.7061039      1.575413  0.002493447      2.3143
63
## 10     2024-01-22      0.6841134      1.537223 -0.001510437      2.2727
93
## 11     2024-01-19      0.8224867      1.468245  0.038410644      2.3103
62
## 12     2024-01-18      0.7989812      1.381504  0.034760044      2.2268
75
## 13     2024-01-17      0.8966285      1.205358  0.077978441      2.1503
45
##      Google_Price Nvidia_Price Berkshire_Price Netflix_Price Amazon_Price
## 8      1.574009      3.061210      1.767665      1.0256221      0.7116560
## 9      1.508753      2.945487      1.684862      0.5857639      0.6805454
## 10     1.467477      2.928466      1.614485      0.5316583      0.6351605
## 11     1.482808      2.915798      1.595055      0.5086134      0.6556569
## 12     1.368807      2.730516      1.507546      0.5283185      0.5883117
## 13     1.289792      2.648600      1.456406      0.4867373      0.5227965
##      Meta_Price Index Bitcoin_Lag var2_7day_avg Bitcoin_Price_Change
## 8      1.899251      8      0.7092066      0.8433251      Decreased
## 9      1.823856      9      0.7191500      0.8131520      Decreased
## 10     1.776975     10      0.7061039      0.7845714      Decreased
## 11     1.799867     11      0.6841134      0.7723021      Decreased
## 12     1.699523     12      0.8224867      0.7533331      Decreased
## 13     1.593148     13      0.7989812      0.7623815      Increased
```

```
table(std_train$Bitcoin_Price_Change)
```

```
##
## Decreased Increased
##      412      341
```

```
table(std_test$Bitcoin_Price_Change)
```

```
##
## Decreased Increased
##      133      113
```

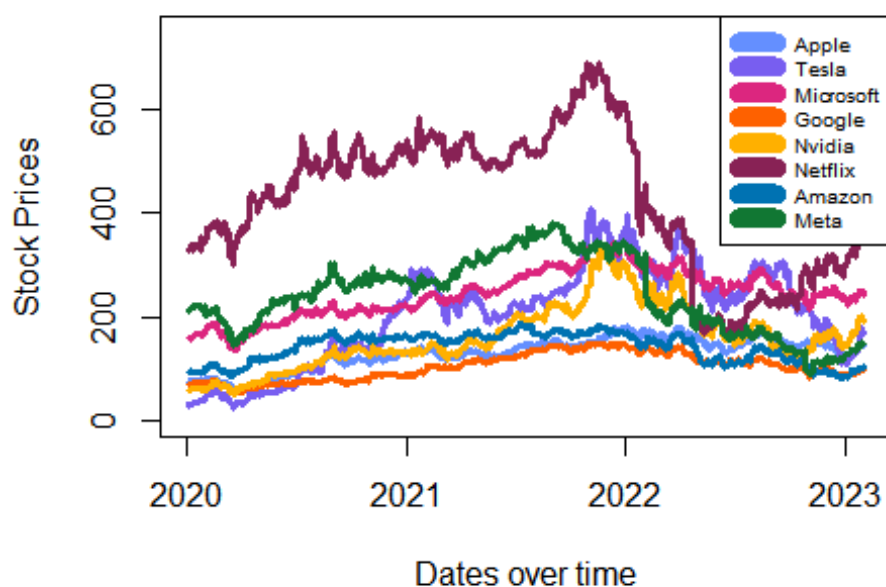

Exploratory Data Analysis

Overlapping Line Graph with Stock Prices

Note: Berkshire_Price was not graphed because it has extreme values

```
plot(train_data$standardized_dates, train_data$Apple_Price, type = "l", col =
"#648fff", lwd = 3, xlab = "Dates over time", ylab = "Stock Prices", main =
"Stock Prices Over Time From Jan 2020 to Jan 2023", ylim = c(0, 750))
lines(train_data$standardized_dates, train_data$Tesla_Price, col = "#785ef0",
lwd = 3)
lines(train_data$standardized_dates, train_data$Microsoft_Price, col = "#dc26
7f", lwd = 3)
lines(train_data$standardized_dates, train_data$Google_Price, col = "#fe6100"
, lwd = 3)
lines(train_data$standardized_dates, train_data$Nvidia_Price, col = "#ffb000"
, lwd = 3)
# lines(train_data$standardized_dates, train_data$Berkshire_Price, col = "#78
5ef0", lwd = 2)
lines(train_data$standardized_dates, train_data$Netflix_Price, col = "#882255
", lwd = 3)
lines(train_data$standardized_dates, train_data$Amazon_Price, col = "#0072b2"
, lwd = 3)
lines(train_data$standardized_dates, train_data$Meta_Price, col = "#117733",
lwd = 3)
legend("topright", legend = c("Apple", "Tesla", "Microsoft", "Google", "Nvidi
a", "Netflix", "Amazon", "Meta"), col = c("#648fff", "#785ef0", "#dc267f", "#f
e6100", "#ffb000", "#882255", "#0072b2", "#117733"), lwd = 10, cex = 0.65)
```

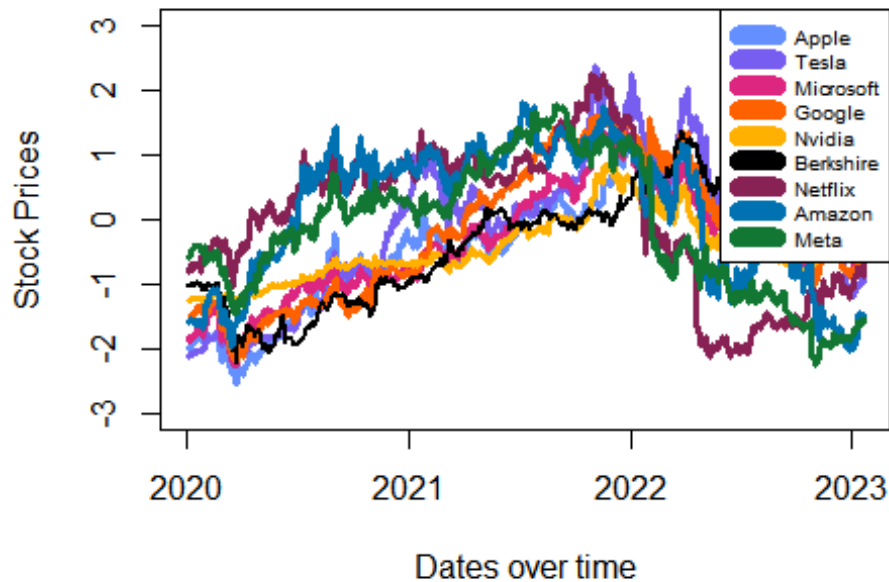
Stock Prices Over Time From Jan 2020 to Jan 202



Overlapping Line Graph with Standardized Stock Prices (Berkshire Included)

```
plot(std_train$standardized_dates, std_train$Apple_Price, type = "l", col = "#648fff", lwd = 3, xlab = "Dates over time", ylab = "Stock Prices", main = "Stock Prices Over Time From Jan 2020 to Jan 2023", ylim = c(-3,3))
lines(std_train$standardized_dates, std_train$Tesla_Price, col = "#785ef0", lwd = 3)
lines(std_train$standardized_dates, std_train$Microsoft_Price, col = "#dc267f", lwd = 3)
lines(std_train$standardized_dates, std_train$Google_Price, col = "#fe6100", lwd = 3)
lines(std_train$standardized_dates, std_train$Nvidia_Price, col = "#ffb000", lwd = 3)
lines(std_train$standardized_dates, std_train$Berkshire_Price, col = "#000000", lwd = 2)
lines(std_train$standardized_dates, std_train$Netflix_Price, col = "#882255", lwd = 3)
lines(std_train$standardized_dates, std_train$Amazon_Price, col = "#0072b2", lwd = 3)
lines(std_train$standardized_dates, std_train$Meta_Price, col = "#117733", lwd = 3)
legend("topright", legend = c("Apple", "Tesla", "Microsoft", "Google", "Nvidia", "Berkshire", "Netflix", "Amazon", "Meta"), col = c("#648fff", "#785ef0", "#dc267f", "#fe6100", "#ffb000", "#000000", "#882255", "#0072b2", "#117733"), lwd = 10, cex = 0.65)
```


Stock Prices Over Time From Jan 2020 to Jan 202



Overlapping Line Graph with Only Apple, Bitcoin, Microsoft, and Tesla

```
# Convert wide data to long format for visualization
std_train_long <- std_train %>%
  pivot_longer(
    cols = c(Bitcoin_Price, Apple_Price, Tesla_Price, Microsoft_Price,
    ),
    names_to = "Stock",
    values_to = "Price"
  )


# Create line plots for all stocks and Bitcoin price over time
ggplot(std_train_long, aes(x = standardized_dates, y = Price, color = Stock))
+
  geom_line() +
  labs(title = "Stock and Bitcoin Price Trends Over Time",
    x = "Date",
    y = "Price",
    color = "Stock") +
  theme_minimal()
```



Overlapping Line Graph with Only Berkshire, Bitcoin, Google, and Nvidia

```
long_data <- std_train %>%
  pivot_longer(cols = c(Bitcoin_Price, Google_Price, Nvidia_Price, Berkshire_Price),
               names_to = "Stock", values_to = "Price")


ggplot(long_data, aes(x = standardized_dates, y = Price, color = Stock)) +
  geom_line() +
  labs(title = "Price Trends Over Time", x = "Date", y = "Price") +
  theme_minimal()
```



Overlapping Line Graph with Only Amazon, Bitcoin, Meta, and Netflix

```
long_data <- std_train %>%
  pivot_longer(cols = c(Bitcoin_Price, Netflix_Price,
                        Amazon_Price, Meta_Price),
               names_to = "Stock", values_to = "Price")

ggplot(long_data, aes(x = standardized_dates, y = Price, color = Stock)) +
  geom_line() +
  labs(title = "Price Trends Over Time", x = "Date", y = "Price") +
  theme_minimal()
```



Correlation Heat Map of all Variables

#Correlation Analysis

#Explore the correlation between Bitcoin's price and the prices of individual stocks


```
correlation_matrix <- cor(train_data[, 2:11], use = "complete.obs")
print(correlation_matrix)
```

```
##          Bitcoin_Price Apple_Price Tesla_Price Microsoft_Price
## Bitcoin_Price      1.0000000  0.62562027   0.7705045    0.7776901
## Apple_Price       0.6256203  1.00000000   0.9188914    0.9101306
## Tesla_Price      0.7705045  0.91889138   1.0000000    0.9016140
## Microsoft_Price   0.7776901  0.91013058   0.9016140    1.0000000
## Google_Price      0.8365944  0.85461788   0.8741953    0.9681007
## Nvidia_Price      0.7300287  0.85800207   0.8527103    0.9589585
## Berkshire_Price   0.5825099  0.86000432   0.7465021    0.7914660
## Netflix_Price     0.4747068 -0.02108043   0.1363454    0.2226306
## Amazon_Price      0.6216082  0.39447118   0.5185501    0.5322344
## Meta_Price        0.5926030  0.16594103   0.3012957    0.4122428
##
##          Google_Price Nvidia_Price Berkshire_Price Netflix_Price
## Bitcoin_Price      0.8365944   0.7300287    0.58250986   0.47470678
## Apple_Price        0.8546179   0.8580021    0.86000432  -0.02108043
## Tesla_Price        0.8741953   0.8527103    0.74650206   0.13634543
## Microsoft_Price    0.9681007   0.9589585    0.79146605   0.22263056
## Google_Price       1.0000000   0.9188245    0.79049291   0.24303381
## Nvidia_Price       0.9188245   1.0000000    0.75083653   0.28381916
```

```
## Berkshire_Price    0.7904929    0.7508365    1.00000000    -0.23557635
## Netflix_Price      0.2430338    0.2838192    -0.23557635    1.00000000
## Amazon_Price       0.5436828    0.5222250    0.08161363    0.76643792
## Meta_Price         0.4832585    0.4146457    -0.02753685    0.84721147
##
##      Amazon_Price  Meta_Price
## Bitcoin_Price    0.62160815  0.59260295
## Apple_Price      0.39447118  0.16594103
## Tesla_Price     0.51855006  0.30129573
## Microsoft_Price  0.53223444  0.41224280
## Google_Price     0.54368278  0.48325853
## Nvidia_Price     0.52222499  0.41464569
## Berkshire_Price  0.08161363 -0.02753685
## Netflix_Price    0.76643792  0.84721147
## Amazon_Price     1.00000000  0.86083900
## Meta_Price       0.86083900  1.00000000
```

Heatmap of correlations

```
ggcorrplot(correlation_matrix, lab = TRUE)
```




#Seasonality or Temporal Trends

#Check if there are any temporal patterns in Bitcoin price changes

```
ggplot(std_train, aes(x = month(standardized_dates), fill = Bitcoin_Price_Change)) +
  geom_bar() +
  labs(title = "Bitcoin Price Change by Month", x = "Month", y = "Frequency")
```

```
+  
theme_minimal()
```



Model building

Logistic Regression Model Based on Stock Prices

Accuracy: $(39 + 81)/246 = 0.4878$ Error Rate: $1 - 0.4878 = 0.5122$ Sensitivity: $81/113 = 0.7168$ Specificity: $39/133 = 0.3451$ Precision: $81/175 = 0.4629$ F1: $(2)(0.4629)(0.7168)/(0.4629 + 0.7168) = 0.5625$ F2: $(5)(0.4629)(0.7168)/((4)(0.4629) + 0.7168) = 0.6459$ F0.5: $(1.25)(0.4629)(0.7168)/((0.25)(0.4629) + 0.7168) = 0.4982$

```
lrs01_model <- glm(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price  
+ Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_  
Price + Amazon_Price + Meta_Price, data = std_train, family = binomial)  
summary(lrs01_model)
```

```
##  
## Call:  
## glm(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price +  
##     Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price +  
##     Netflix_Price + Amazon_Price + Meta_Price, family = binomial,  
##     data = std_train)  
##  
## Coefficients:
```



```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.3216    0.2169   1.483 0.138133
## Apple_Price   -1.8567    0.4207  -4.413 1.02e-05 ***
## Tesla_Price    0.8394    0.2234   3.758 0.000172 ***
## Microsoft_Price 1.6823    0.7121   2.362 0.018156 *
## Google_Price  -1.4646    0.4653  -3.148 0.001646 **
## Nvidia_Price   0.8578    0.6194   1.385 0.166051
## Berkshire_Price 0.8860    0.2878   3.079 0.002080 **
## Netflix_Price  -0.6075    0.1752  -3.468 0.000524 ***
## Amazon_Price   0.2561    0.2197   1.166 0.243669
## Meta_Price     0.4536    0.2542   1.785 0.074317 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 1037.18  on 752  degrees of freedom
## Residual deviance:  972.61  on 743  degrees of freedom
## AIC: 992.61
##
## Number of Fisher Scoring iterations: 4

print("Predictor Variable Multicollinearity:")

## [1] "Predictor Variable Multicollinearity:"

ypred <- predict(object = lrs01_model, newdata = std_test, type = "response")
predicted_classes <- ifelse(ypred > 0.7, 1, 0)

# Contingency Table

t1 <- table(std_test$Bitcoin_Price_Change, predicted_classes)
row.names(t1) <- c("Actual: Decrease", "Actual: Increase")
colnames(t1) <- c("Predicted: Decrease", "Predicted: Increase")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1

##              predicted_classes
##              Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                39                 94    133
## Actual: Increase                32                 81    113
## Total                          71                175    246

# Logistic Regression AUC score

lrs01_actual <- (std_test$Bitcoin_Price_Change == "Increased")
lrs01_roc <- roc(lrs01_actual, predicted_classes)

## Setting levels: control = FALSE, case = TRUE

## Setting direction: controls < cases
```

```
lrs01_auc <- auc(lrs01_roc)
lrs01_auc

## Area under the curve: 0.505
```

Logistic Regression Model with Removed Predictor Variables

Accuracy: $(23 + 100)/246 = 0.5000$ Error Rate: $1 - 0.5000 = 0.5000$ Sensitivity: $100/113 = 0.8850$ Specificity: $23/133 = 0.1729$ Precision: $100/210 = 0.4762$ F1: $(2)(0.4762)(0.8850)/(0.4762 + 0.8850) = 0.6192$ F2: $(5)(0.4762)(0.8850)/((4)(0.4762) + 0.8850) = 0.7553$ F0.5: $(1.25)(0.4762)(0.8850)/((0.25)(0.4762) + 0.8850) = 0.5247$

```
lrs02_model <- glm(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price
+ Microsoft_Price + Google_Price + Berkshire_Price, data = std_train, family
= binomial)
summary(lrs02_model)

##
## Call:
## glm(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price +
##     Microsoft_Price + Google_Price + Berkshire_Price, family = binomial,
##     data = std_train)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.0001739  0.1292146  -0.001  0.998926
## Apple_Price   -1.4123601  0.3714315  -3.802  0.000143 ***
## Tesla_Price    0.7366876  0.2107599   3.495  0.000473 ***
## Microsoft_Price 1.4869419  0.5271918   2.820  0.004795 **
## Google_Price  -0.8823492  0.3499998  -2.521  0.011702 *
## Berkshire_Price 0.7415763  0.2038872   3.637  0.000276 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 1037.2  on 752  degrees of freedom
## Residual deviance:  985.9  on 747  degrees of freedom
## AIC: 997.9
##
## Number of Fisher Scoring iterations: 4

print("Predictor Variable Multicollinearity:")

## [1] "Predictor Variable Multicollinearity:"

vif(lrs02_model)

##      Apple_Price      Tesla_Price Microsoft_Price      Google_Price Berkshire_
Price
```

```
##          18.825248          9.230368          30.850198          21.932795          5.0
49894

ypred <- predict(object = lrs02_model, newdata = std_test, type = "response")
predicted_classes <- ifelse(ypred > 0.5, 1, 0)

# Contingency Table

t2 <- table(std_test$Bitcoin_Price_Change, predicted_classes)
row.names(t2) <- c("Actual: Decrease", "Actual: Increase")
colnames(t2) <- c("Predicted: Decrease", "Predicted: Increase")
t2 <- addmargins(A = t2, FUN = list(Total = sum), quiet = TRUE)
t2

##                predicted_classes
##                Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                23                110    133
## Actual: Increase                13                100    113
## Total                          36                210    246

# Logistic Regression AUC score

lrs02_actual <- (std_test$Bitcoin_Price_Change == "Increased")
lrs02_roc <- roc(lrs02_actual, predicted_classes)

## Setting levels: control = FALSE, case = TRUE

## Setting direction: controls < cases

lrs02_auc <- auc(lrs02_roc)
lrs02_auc

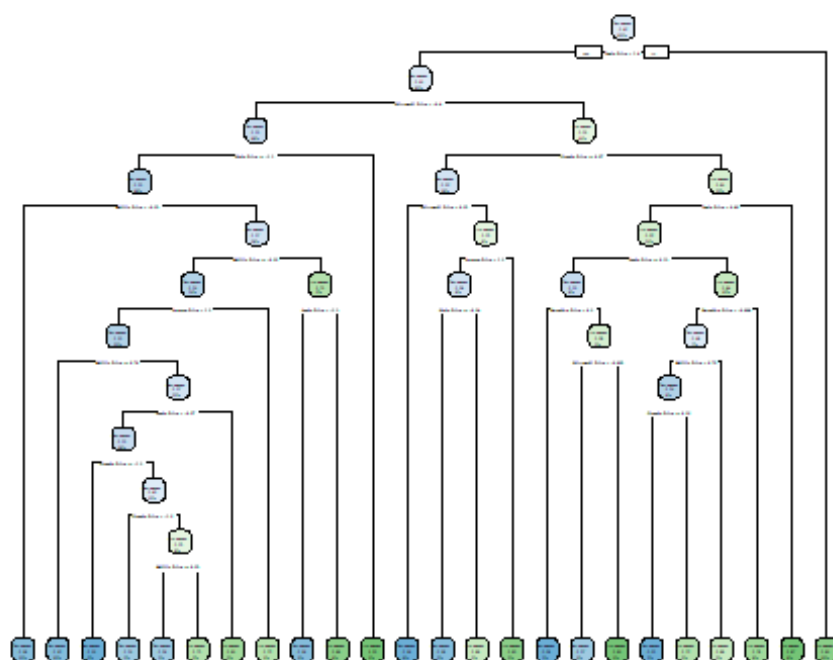
## Area under the curve: 0.5289
```

CART

Accuracy: $(83 + 55)/246 = 0.5610$ Error Rate: $1 - 0.5610 = 0.4390$ Sensitivity: $55/113 = 0.4867$ Specificity: $83/133 = 0.6241$ Precision: $55/105 = 0.5238$ F1: $(2)(0.5238)(0.4867)/(0.5238 + 0.4867) = 0.5046$ F2: $(5)(0.5238)(0.4867)/((4)(0.5238) + 0.4867) = 0.4937$ F0.5: $(1.25)(0.5238)(0.4867)/((0.25)(0.5238) + 0.4867) = 0.5159$

```
cart01_model <- rpart(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Pr
ice + Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netfl
ix_Price + Amazon_Price + Meta_Price, data = std_train, method = "class")

rpart.plot(cart01_model)
```



```
ypred <- predict(object = cart01_model, newdata = std_test, type = "class")
# predicted_class <- ifelse(ypred[, 2] > 0.5, 1, 0)
```

Contingency Table

```
t1 <- table(std_test$Bitcoin_Price_Change, ypred)
row.names(t1) <- c("Actual: Decrease", "Actual: Increase")
colnames(t1) <- c("Predicted: Decrease", "Predicted: Increase")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1
```

```
##               ypred
##               Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                83                 50   133
## Actual: Increase                58                 55   113
## Total                          141                105   246
```

Node counts

```
total_nodes <- length(cart01_model$frame$var)
total_nodes
```

```
## [1] 47
```

```
leaf_nodes <- sum(cart01_model$frame$var == "<leaf>")
leaf_nodes
```



```

t2 <- table(std_test$Bitcoin_Price_Change, predicted_classes)
row.names(t2) <- c("Actual: Decrease", "Actual: Increase")
colnames(t2) <- c("Predicted: Decrease", "Predicted: Increase")
t2 <- addmargins(A = t2, FUN = list(Total = sum), quiet = TRUE)
t2

##               predicted_classes
##               Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                74                 59    133
## Actual: Increase                49                 64    113
## Total                          123                123    246

# Node counts

total_nodes <- length(cart02_model$frame$var)
total_nodes

## [1] 29

leaf_nodes <- sum(cart02_model$frame$var == "<leaf>")
leaf_nodes

## [1] 15

decision_nodes <- total_nodes - leaf_nodes
decision_nodes

## [1] 14

# CART AUC score

cart02_actual <- (std_test$Bitcoin_Price_Change == "Increased")
cart02_roc <- roc(cart02_actual, predicted_classes)

## Setting levels: control = FALSE, case = TRUE

## Setting direction: controls < cases

cart02_auc <- auc(cart02_roc)
cart02_auc

## Area under the curve: 0.5614

```

Random Forests

Accuracy: $(41 + 92)/246 = 0.5407$ Error Rate: $1 - 0.4919 = 0.4593$ Sensitivity: $92/113 = 0.8142$ Specificity: $41/133 = 0.3083$ Precision: $92/184 = 0.4457$ F1: $(2)(0.4457)(0.8142)/(0.4457 + 0.8142) = 0.5761$ F2: $(5)(0.4457)(0.8142)/((4)(0.4457) + 0.8142) = 0.6987$ F0.5: $(1.25)(0.4457)(0.8142)/((0.25)(0.4457) + 0.8142) = 0.4901$

```

set.seed(35)
rf01_model <- randomForest(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price + Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_Price + Amazon_Price + Meta_Price, data = std_train, method = "class")

ypred <- predict(object = rf01_model, newdata = std_test, type = "class")

# Contingency Table

t1 <- table(std_test$Bitcoin_Price_Change, ypred)
row.names(t1) <- c("Actual: Decrease", "Actual: Increase")
colnames(t1) <- c("Predicted: Decrease", "Predicted: Increase")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1

##                ypred
##                Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                41                92    133
## Actual: Increase                21                92    113
## Total                          62               184    246

rf01_model$ntree

## [1] 500

# RF AUC score

rf01_roc <- roc(std_test$Bitcoin_Price_Change, as.numeric(ypred))

## Setting levels: control = Decreased, case = Increased
## Setting direction: controls < cases

rf01_auc <- auc(rf01_roc)
rf01_auc

## Area under the curve: 0.5612

```

Random Forests with 1000 ntree

```

set.seed(35)
rf02_model <- randomForest(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price + Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_Price + Amazon_Price + Meta_Price, data = std_train, method = "class", ntree = 1000)

ypred <- predict(object = rf02_model, newdata = std_test, type = "class")

# Contingency Table

```

```
t2 <- table(std_test$Bitcoin_Price_Change, ypred)
row.names(t2) <- c("Actual: Decrease", "Actual: Increase")
colnames(t2) <- c("Predicted: Decrease", "Predicted: Increase")
t2 <- addmargins(A = t2, FUN = list(Total = sum), quiet = TRUE)
t2

##                ypred
##                Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                64                69    133
## Actual: Increase                40                73    113
## Total                          104               142    246

rf02_model$ntree

## [1] 1000

# RF AUC score

rf02_roc <- roc(std_test$Bitcoin_Price_Change, as.numeric(ypred))

## Setting levels: control = Decreased, case = Increased

## Setting direction: controls < cases

rf02_auc <- auc(rf02_roc)
rf02_auc

## Area under the curve: 0.5636
```

Random Forests with 100 ntree

```
set.seed(35)
rf03_model <- randomForest(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price + Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_Price + Amazon_Price + Meta_Price, data = std_train, method = "class", ntree = 100)

ypred <- predict(object = rf03_model, newdata = std_test, type = "class")

# Contingency Table

t3 <- table(std_test$Bitcoin_Price_Change, ypred)
row.names(t3) <- c("Actual: Decrease", "Actual: Increase")
colnames(t3) <- c("Predicted: Decrease", "Predicted: Increase")
t3 <- addmargins(A = t3, FUN = list(Total = sum), quiet = TRUE)
t3
```

	ypred			
		Predicted: Decrease	Predicted: Increase	Total
##	Actual: Decrease	37	96	133
##	Actual: Increase	27	86	113
##	Total	64	182	246


```

rf03_model$ntree
## [1] 100
# RF AUC score

rf03_roc <- roc(std_test$Bitcoin_Price_Change, as.numeric(ypred))
## Setting levels: control = Decreased, case = Increased
## Setting direction: controls < cases

rf03_auc <- auc(rf03_roc)
rf03_auc

## Area under the curve: 0.5196

```

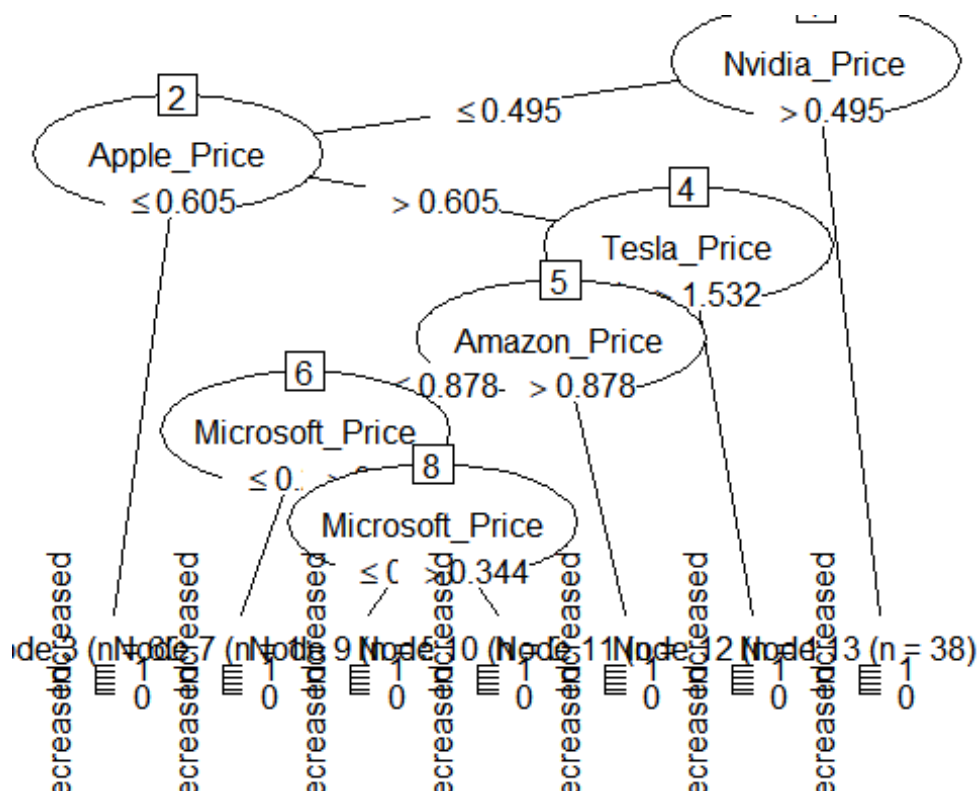
C5.0

```

c5_model <- C5.0(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price +
  Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_Pr
ice + Amazon_Price + Meta_Price, data = std_train, methond = "class")

plot(c5_model)

```



```

ypred <- predict(object = c5_model, newdata = std_test, type = "class")

t1 <- table(std_test$Bitcoin_Price_Change, ypred)

```

```

row.names(t1) <- c("Actual: Decrease", "Actual: Increase")
colnames(t1) <- c("Predicted: Decrease", "Predicted: Increase")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1

##                ypred
##                Predicted: Decrease Predicted: Increase Total
## Actual: Decrease                26                107    133
## Actual: Increase                19                 94    113
## Total                          45                201    246

```

Neural Networks


```

set.seed(35)
nnet01_model <- nnet(formula = Bitcoin_Price_Change ~ Apple_Price + Tesla_Price +
                    Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price +
                    Netflix_Price + Amazon_Price + Meta_Price, data = std_train, size = 1)

## # weights: 12
## initial value 518.573865
## iter 10 value 486.752168
## iter 20 value 462.476803
## iter 30 value 458.052086
## iter 40 value 457.592883
## iter 50 value 456.919594
## iter 60 value 456.493738
## iter 70 value 456.328923
## iter 80 value 456.305374
## iter 90 value 456.275022
## iter 100 value 456.262995
## final value 456.262995
## stopped after 100 iterations

plotnet(nnet01_model)

```



```
ypred <- predict(object = nnet01_model, newdata = std_test, type = "class")
```

```
t1 <- table(std_test$Bitcoin_Price_Change, ypred)
row.names(t1) <- c("Actual: Decrease", "Actual: Increase")
colnames(t1) <- c("Predicted: Decrease", "Predicted: Increase")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1
```

	ypred		
	Predicted: Decrease	Predicted: Increase	Total
Actual: Decrease	40	93	133
Actual: Increase	45	68	113
Total	85	161	246

Naive Bayes

```
# Train the Naive Bayes model
nb_model <- naiveBayes(Bitcoin_Price_Change ~ Apple_Price + Tesla_Price + Microsoft_Price + Google_Price + Nvidia_Price + Berkshire_Price + Netflix_Price + Amazon_Price + Meta_Price, data = std_train)
```

```
# Predict on the test set
nb_predictions_prob <- predict(nb_model, newdata = std_test, type = "raw")
nb_predictions <- ifelse(nb_predictions_prob[, "Increased"] >= 0.9, "Increased", "Decreased")
```

```
# Confusion matrix
```

```

conf_matrix <- table(std_test$Bitcoin_Price_Change, nb_predictions)

# Add proper row and column names
row.names(conf_matrix) <- c("Actual: Decrease", "Actual: Increase")
colnames(conf_matrix) <- c("Predicted: Decrease", "Predicted: Increase")

# Print the confusion matrix
print(conf_matrix)

##              nb_predictions
##              Predicted: Decrease Predicted: Increase
## Actual: Decrease              35              98
## Actual: Increase              27              86

```